

2j8x112qd

March 3, 2024

```
[1]: !pip install prophet
```

```
Requirement already satisfied: prophet in c:\users\vince\anaconda3\lib\site-  
packages (1.1.5)  
Requirement already satisfied: cmdstanpy>=1.0.4 in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (1.2.1)  
Requirement already satisfied: numpy>=1.15.4 in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (1.24.3)  
Requirement already satisfied: matplotlib>=2.0.0 in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (3.7.1)  
Requirement already satisfied: pandas>=1.0.4 in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (1.5.3)  
Requirement already satisfied: holidays>=0.25 in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (0.43)  
Requirement already satisfied: tqdm>=4.36.1 in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (4.65.0)  
Requirement already satisfied: importlib-resources in  
c:\users\vince\anaconda3\lib\site-packages (from prophet) (6.1.2)  
Requirement already satisfied: stanio~=0.3.0 in  
c:\users\vince\anaconda3\lib\site-packages (from cmdstanpy>=1.0.4->prophet)  
(0.3.0)  
Requirement already satisfied: python-dateutil in  
c:\users\vince\anaconda3\lib\site-packages (from holidays>=0.25->prophet)  
(2.8.2)  
Requirement already satisfied: contourpy>=1.0.1 in  
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)  
(1.0.5)  
Requirement already satisfied: cycler>=0.10 in  
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)  
(0.11.0)  
Requirement already satisfied: fonttools>=4.22.0 in  
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)  
(4.25.0)  
Requirement already satisfied: kiwisolver>=1.0.1 in  
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)  
(1.4.4)  
Requirement already satisfied: packaging>=20.0 in  
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)
```

(23.0)
Requirement already satisfied: pillow>=6.2.0 in
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)
(9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
c:\users\vince\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet)
(3.0.9)
Requirement already satisfied: pytz>=2020.1 in
c:\users\vince\anaconda3\lib\site-packages (from pandas>=1.0.4->prophet)
(2022.7)
Requirement already satisfied: colorama in c:\users\vince\anaconda3\lib\site-
packages (from tqdm>=4.36.1->prophet) (0.4.6)
Requirement already satisfied: six>=1.5 in c:\users\vince\anaconda3\lib\site-
packages (from python-dateutil->holidays>=0.25->prophet) (1.16.0)

```
[2]: # Import Libraries and packages

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("whitegrid")
import datetime
import math
import warnings
import prophet
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.stattools import adfuller
from tensorflow.keras.layers import LSTM, Dense
from tensorflow.keras.models import Sequential

warnings.filterwarnings("ignore")
```

WARNING:tensorflow:From C:\Users\vince\anaconda3\Lib\site-
packages\keras\src\losses.py:2976: The name
tf.losses.sparse_softmax_cross_entropy is deprecated. Please use
tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
[3]: # import and display csv file
df = pd.read_csv('AAPL.csv')
df.head()
```

```
[3]:
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	1980-12-12	0.128348	0.128906	0.128348	0.128348	0.100323	469033600
1	1980-12-15	0.122210	0.122210	0.121652	0.121652	0.095089	175884800
2	1980-12-16	0.113281	0.113281	0.112723	0.112723	0.088110	105728000
3	1980-12-17	0.115513	0.116071	0.115513	0.115513	0.090291	86441600
4	1980-12-18	0.118862	0.119420	0.118862	0.118862	0.092908	73449600

```
[4]: # display bottom rows of csv file
df.tail()
```

```
[4]:
```

	Date	Open	High	Low	Close	Adj Close	\
10404	2022-03-18	160.509995	164.479996	159.759995	163.979996	163.979996	
10405	2022-03-21	163.509995	166.350006	163.009995	165.380005	165.380005	
10406	2022-03-22	165.509995	169.419998	164.910004	168.820007	168.820007	
10407	2022-03-23	167.990005	172.639999	167.649994	170.210007	170.210007	
10408	2022-03-24	171.059998	174.139999	170.210007	174.070007	174.070007	

	Volume
10404	123351200
10405	95811400
10406	81532000
10407	98062700
10408	90018700

```
[5]: # checking shape of dataframe
df.shape
```

```
[5]: (10409, 7)
```

```
[6]: #checking data types of data frame
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10409 entries, 0 to 10408
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        10409 non-null  object
1   Open        10409 non-null  float64
2   High        10409 non-null  float64
3   Low         10409 non-null  float64
4   Close       10409 non-null  float64
5   Adj Close   10409 non-null  float64
6   Volume      10409 non-null  int64
dtypes: float64(5), int64(1), object(1)
memory usage: 569.4+ KB
```

```
[7]: #renaming columns to python casing.
col_head = {
    'Date': 'date',
    'Open': 'open',
    'High': 'high',
    'Close': 'close',
    'Low': 'low',
    'Adj Close': 'adj_close',
    'Volume': 'volume'}
df.rename(columns=col_head, inplace=True)
```

```
[8]: # checking for NULL values
print("Null Values:\n", df.isna().sum())

# dropping null values
df = df.dropna()

# verifying null values were dropped
print("Null Values after dropping:\n", df.isna().sum())
```

Null Values:

date	0
open	0
high	0
low	0
close	0
adj_close	0
volume	0

dtype: int64

Null Values after dropping:

date	0
open	0
high	0
low	0
close	0
adj_close	0
volume	0

dtype: int64

```
[9]: # checking for duplicate values
print("Duplicate Values:\n", df.duplicated().sum())

# dropping duplicate values
df = df.drop_duplicates()

# verifying duplicates were dropped
print("Duplicate Values after dropping:\n", df.duplicated().sum())
```

Duplicate Values:

0

Duplicate Values after dropping:







0

```
[10]: # Create a figure and a grid of subplots
fig, axs = plt.subplots(nrows=6, ncols=1, figsize=(10, 20))

# plot each trace on the corresponding subplot
axs[0].plot(df.index, df["open"])
axs[0].set_title("Opening Price")
axs[1].plot(df.index, df["close"])
axs[1].set_title("Closing Price")
axs[2].plot(df.index, df["high"])
axs[2].set_title("Highest Price")
axs[3].plot(df.index, df["low"])
axs[3].set_title("Lowest Price")
axs[4].plot(df.index, df["adj_close"])
axs[4].set_title("Adjusted Closing Price")
axs[5].plot(df.index, df["volume"])
axs[5].set_title("Volume")

# Adjust layout
plt.tight_layout()

# Show the plot
```



```
[11]: # printing summary statistics
df.describe()
```

```
[11]:
```

	open	high	low	close	adj_close \
count	10409.000000	10409.000000	10409.000000	10409.000000	10409.000000
mean	13.959910	14.111936	13.809163	13.966757	13.350337
std	30.169244	30.514878	29.835055	30.191696	29.911132
min	0.049665	0.049665	0.049107	0.049107	0.038384
25%	0.281964	0.287946	0.274554	0.281250	0.234799
50%	0.468750	0.477679	0.459821	0.468750	0.386853
75%	14.217857	14.364286	14.043571	14.206071	12.188149
max	182.630005	182.940002	179.119995	182.009995	181.778397

	volume
count	1.040900e+04
mean	3.321778e+08
std	3.393344e+08
min	0.000000e+00
25%	1.247604e+08
50%	2.199680e+08
75%	4.126108e+08
max	7.421641e+09

```
[12]: # importing plotly for an interactive graph
import plotly.graph_objects as go
import plotly.express as px

# create a figure
fig = px.line(df, x=df.index, y=['open', 'high', 'low', 'close'], title='Apple_
↳ Stock Price 1980-2022')

# add labels
fig.update_layout(xaxis_title='Date', yaxis_title='Price (USD)')

# show the interactive plot
fig.show()
```

```
[13]: # calculate the moving average
moving_avg = df['close'].rolling(window=100).mean()

# create a figure
fig = go.Figure()

# add the moving average trace
```

```

fig.add_trace(go.Scatter(x=df.index, y=moving_avg, mode='lines', name='Simple_
↳Moving Average'))

# add the closing price trace
fig.add_trace(go.Scatter(x=df.index, y=df["close"], mode='lines', name='Closing_
↳Price'))

# update layout
fig.update_layout(title='Moving Average of Closing Price',
                  xaxis_title='Date',
                  yaxis_title='Price (USD)',
                  xaxis=dict(tickangle=-45),
                  showlegend=True)

# show the plot
fig.show()

```

```

[14]: # create subset data frame that contains just the date and close column
df_close = df[['date', 'close']]
df_close.head()

```

```

[14]:
      date      close
0  1980-12-12  0.128348
1  1980-12-15  0.121652
2  1980-12-16  0.112723
3  1980-12-17  0.115513
4  1980-12-18  0.118862

```

```

[15]: # changing datatype of date column
df_close["date"] = pd.to_datetime(pd.to_datetime(df["date"]).dt.date)
df_close.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10409 entries, 0 to 10408
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
0   date    10409 non-null   datetime64[ns]
1   close   10409 non-null   float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 244.0 KB

```

```

[16]: # updating column names to the Prophet casing convention
df_new = df_close.rename(mapper={"date": "ds", "close": "y"}, axis="columns")
df_new.head()

```




```
[16]:      ds      y
      0 1980-12-12  0.128348
      1 1980-12-15  0.121652
      2 1980-12-16  0.112723
      3 1980-12-17  0.115513
      4 1980-12-18  0.118862
```

```
[17]: # training and forecasting the model using Facebook prophet model
op = prophet.Prophet()
op.fit(df_new)
forecast = op.make_future_dataframe(periods=365)
forecast = op.predict(forecast)
```

17:00:34 - cmdstanpy - INFO - Chain [1] start processing

17:00:39 - cmdstanpy - INFO - Chain [1] done processing

```
[18]: # visualizing forecast results
op.plot(forecast, xlabel="Date", ylabel="Stock End Price")
plt.show()
```



```
[19]: # first forecast was poor, removing data prior to 2015
df_new['ds'] = pd.to_datetime(df_new['ds'])
df_new = df_new[df_new['ds'].dt.year >= 2015]
```

```
[20]: df_new
```

```
[20]:
```

	ds	y
8589	2015-01-02	27.332500
8590	2015-01-05	26.562500
8591	2015-01-06	26.565001
8592	2015-01-07	26.937500
8593	2015-01-08	27.972500
...
10404	2022-03-18	163.979996
10405	2022-03-21	165.380005
10406	2022-03-22	168.820007
10407	2022-03-23	170.210007
10408	2022-03-24	174.070007


[1820 rows x 2 columns]

```
[21]: # training and forecasting the model using Facebook prophet model
op = prophet.Prophet()
op.fit(df_new)
forecast = op.make_future_dataframe(periods=365)
forecast = op.predict(forecast)
```

17:00:47 - cmdstanpy - INFO - Chain [1] start processing

17:00:48 - cmdstanpy - INFO - Chain [1] done processing

```
[22]: # visualizing second forecast results
op.plot(forecast, xlabel="Date", ylabel="Stock End Price")
plt.show()
```




```
[23]: # import seasonal decompose from statsmodels library
from statsmodels.tsa.seasonal import seasonal_decompose

# index on date
df_new.set_index("ds", inplace=True)

# Decomposition
result = seasonal_decompose(df_new.y, model='multiplicative', period=225)

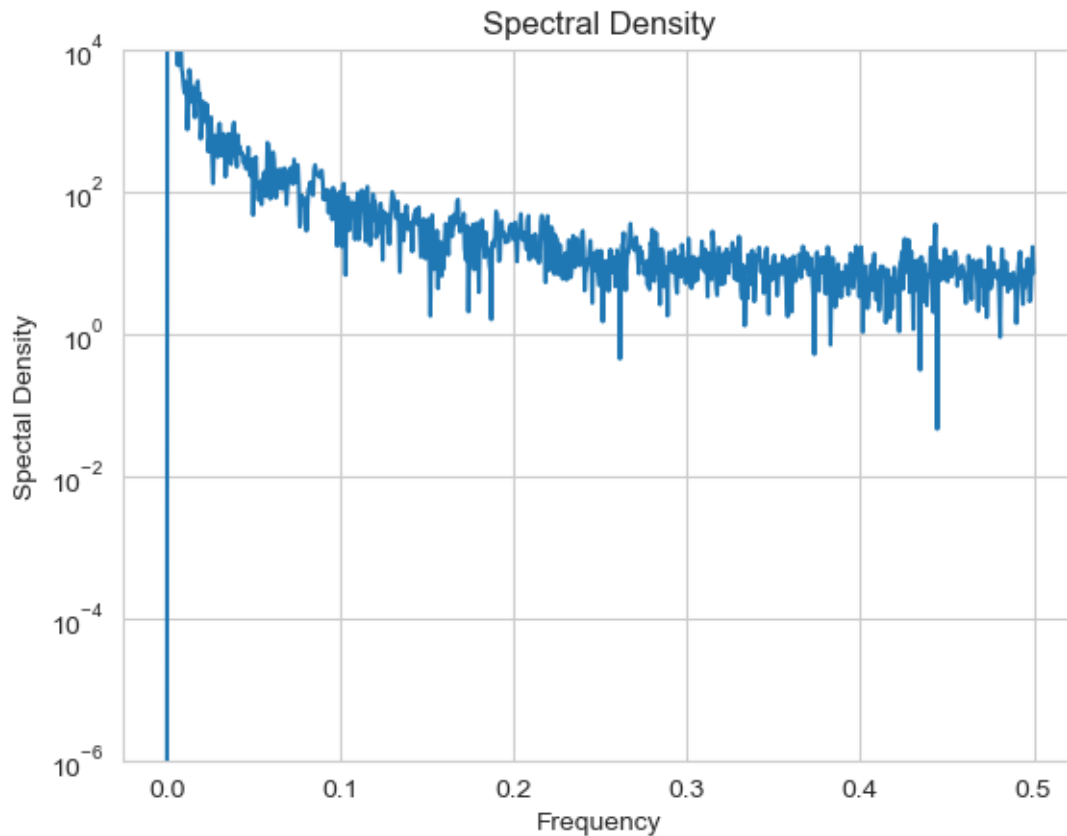
# Plot Trend, Seasonality, Residuals
result.plot()
plt.show()
```



```
[24]: import matplotlib.gridspec as gridspec

# Power spectral density
from scipy import signal
f, Pxx_den = signal.periodogram(df_new['y'])
```

```
plt.semilogy(f, Pxx_den)
plt.ylim([1e-6, 1e4])
plt.title('Spectral Density')
plt.xlabel('Frequency')
plt.ylabel('Spectral Density')
plt.show()
```



```
[39]: # first forecast was poor, removing data prior to 2015
df_close = df_close[df_close['date'].dt.year >= 2015]

# printing the shape of the new model
df_close.shape
```

```
[39]: (1820, 2)
```

```
[40]: df_close.head()
```

```
[40]:
```

	date	close
8589	2015-01-02	27.332500
8590	2015-01-05	26.562500

```
8591 2015-01-06 26.565001
8592 2015-01-07 26.937500
8593 2015-01-08 27.972500
```

```
[41]: df_close.tail()
```

```
[41]:      date      close
10404 2022-03-18 163.979996
10405 2022-03-21 165.380005
10406 2022-03-22 168.820007
10407 2022-03-23 170.210007
10408 2022-03-24 174.070007
```

```
[42]: # splitting the data into training and testing sets 80/20 split
X_train = df_close[:1456]
X_test = df_close[1456:]

# print the shape of the datasets
print("X_train Shape", X_train.shape)
print("X_test Shape", X_test.shape)
```

```
X_train Shape (1456, 2)
X_test Shape (364, 2)
```

```
[43]: # indexing dataset on date
X_train = X_train[["date", "close"]]
X_test = X_test[["date", "close"]]

X_train.set_index("date", inplace=True)
X_test.set_index("date", inplace=True)
```

```
[44]: # creating forecast for next twelve months
fy_forecast = pd.date_range(X_test.index[-1], freq="MS", periods=12)
fy_forecast
```

```
[44]: DatetimeIndex(['2022-04-01', '2022-05-01', '2022-06-01', '2022-07-01',
                    '2022-08-01', '2022-09-01', '2022-10-01', '2022-11-01',
                    '2022-12-01', '2023-01-01', '2023-02-01', '2023-03-01'],
                    dtype='datetime64[ns]', freq='MS')
```

```
[45]: # performing ADF test to check for stationarity
def define_d(data):
    adf_result = adfuller(data, autolag="AIC")
    adf_statistic = adf_result[0]
    pvalue = adf_result[1]
    critical_value = adf_result[4]["5%"]
```

```

print(f"ADF Test Statistic: {adf_statistic}")
print(f"p-value: {pvalue}")
print(f"Critical Value (5%): {critical_value}")

if pvalue < 0.05:
    print("** Data is stationary, proceed to plotting ACF and PACF.**")
else:
    print("** Data is not stationary, and needs to be differenced! **")

define_d(X_train)

```

```

ADF Test Statistic: 2.3635665531945196
p-value: 0.9989921192624683
Critical Value (5%): -2.863561857668172
** Data is not stationary, and needs to be differenced! **

```

```

[46]: # performing differencing to make data stationary
data_diff = X_train.diff()
data_diff.dropna(inplace=True)
define_d(data_diff)

```

```

ADF Test Statistic: -7.488965986870958
p-value: 4.5589820085980184e-11
Critical Value (5%): -2.863561857668172
** Data is stationary, proceed to plotting ACF and PACF.**

```

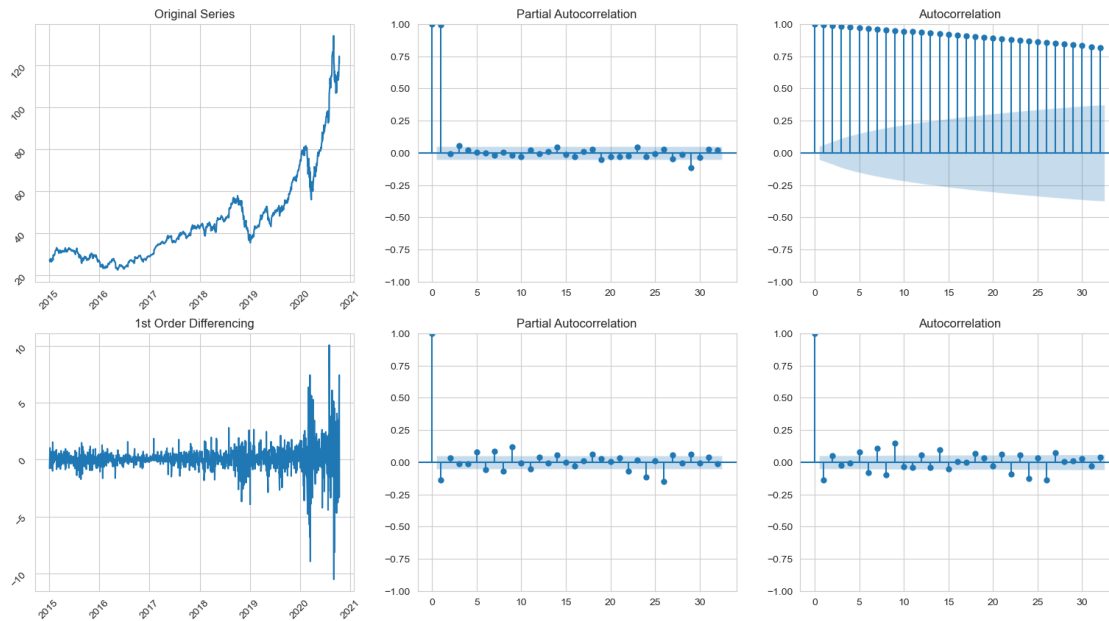
```

[47]: # Original dataframe visualization
fig, axes = plt.subplots(2, 3)
fig.set_figheight(10)
fig.set_figwidth(19)
axes[0, 0].plot(X_train)
axes[0, 0].tick_params(labelrotation=45)
axes[0, 0].set_title("Original Series")
plot_pacf(X_train, ax=axes[0, 1])
plot_acf(X_train, ax=axes[0, 2])

# visualizing data post first differencing
axes[1, 0].plot(X_train.diff())
axes[1, 0].set_title("1st Order Differencing")
axes[1, 0].tick_params(labelrotation=45)
plot_pacf(X_train.diff().dropna(), ax=axes[1, 1])
plot_acf(X_train.diff().dropna(), ax=axes[1, 2])

plt.show();

```



```
[49]: price_validate = df_close["close"][1456:]
price_validate
```

```
[49]: 10045    121.190002
      10046    120.709999
      10047    119.019997
      10048    115.980003
      10049    117.510002
      ...
      10404    163.979996
      10405    165.380005
      10406    168.820007
      10407    170.210007
      10408    174.070007
      Name: close, Length: 364, dtype: float64
```

```
[51]: # calculating the Root Mean Squared Error (RMSE)
def forecast_accuracy(forecast, actual):
    mape = (np.mean(np.abs(forecast - actual) / np.abs(actual)) * 100).round(2)
    rmse = np.sqrt(((forecast - actual) ** 2).mean())
    return {"Mean Absolute Percentage Error (%)": mape, "Root Mean Squared_
↳Error": rmse}
```

0.0.1 Arima Model

```
[52]: model = ARIMA(X_train)
model = model.fit()
print(model.summary())

fc = model.forecast(364)
forecast_accuracy(fc, price_validate.values)
```

SARIMAX Results

```
=====
Dep. Variable:          close    No. Observations:          1456
Model:                  ARIMA    Log Likelihood             -6503.753
Date:                   Sat, 02 Mar 2024    AIC                13011.507
Time:                   17:20:09    BIC                13022.074
Sample:                 0    HQIC                13015.449
                        - 1456
```

```
Covariance Type:          opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	45.3707	0.857	52.959	0.000	43.692	47.050
sigma2	444.0009	16.194	27.417	0.000	412.260	475.741

```
=====
```

```
====
Ljung-Box (L1) (Q):          1441.09    Jarque-Bera (JB):
1240.27
Prob(Q):                      0.00    Prob(JB):
0.00
Heteroskedasticity (H):        3.13    Skew:
1.70
Prob(H) (two-sided):          0.00    Kurtosis:
5.97
=====
====
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
```

```
[52]: {'Mean Absolute Percentage Error (%)': 67.55,
      'Root Mean Squared Error': 98.41523686774337}
```

```
[53]: model = ARIMA(X_train, order=(1, 0, 1))
model = model.fit()
print(model.summary())

fc = model.forecast(364)
```



```
forecast_accuracy(fc, price_validate.values)
```

SARIMAX Results

```
=====
Dep. Variable:          close    No. Observations:          1456
Model:                ARIMA(1, 0, 1)    Log Likelihood          -2251.763
Date:                 Sat, 02 Mar 2024    AIC                    4511.525
Time:                 17:21:16    BIC                    4532.659
Sample:                0    HQIC                    4519.410
                        - 1456
```

```
Covariance Type:          opg
```

```
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
const         46.2916    193.929      0.239      0.811    -333.802     426.385
ar.L1          0.9998      0.001   1292.451      0.000      0.998      1.001
ma.L1         -0.1219      0.011   -11.080      0.000     -0.143     -0.100
sigma2         1.2840      0.014    90.019      0.000      1.256      1.312
=====
```

```
===
Ljung-Box (L1) (Q):                0.15    Jarque-Bera (JB):
26074.85
Prob(Q):                0.70    Prob(JB):
0.00
Heteroskedasticity (H):            16.45    Skew:
-0.27
Prob(H) (two-sided):            0.00    Kurtosis:
23.72
=====
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
```

```
[53]: {'Mean Absolute Percentage Error (%)': 15.75,
      'Root Mean Squared Error': 30.536918644529457}
```

```
[54]: model = ARIMA(X_train, order=(1, 1, 1))
      model = model.fit()
      print(model.summary())

      fc = model.forecast(364)
      forecast_accuracy(fc, price_validate.values)
```

SARIMAX Results

```
=====
Dep. Variable:          close    No. Observations:          1456
```

```

Model:          ARIMA(1, 1, 1)    Log Likelihood    -2244.252
Date:           Sat, 02 Mar 2024    AIC             4494.504
Time:           17:21:20           BIC             4510.352
Sample:         0                   HQIC             4500.417
                  - 1456

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         -0.3705      0.067     -5.551      0.000     -0.501     -0.240
ma.L1          0.2398      0.070      3.413      0.001      0.102      0.377
sigma2         1.2802      0.014     90.495      0.000      1.252      1.308
=====

```

===

```

Ljung-Box (L1) (Q):          0.00    Jarque-Bera (JB):
26453.30
Prob(Q):          0.94    Prob(JB):
0.00
Heteroskedasticity (H):      16.24    Skew:
-0.25
Prob(H) (two-sided):        0.00    Kurtosis:
23.88

```

```

=====
===

```

Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients (complex-
step).

```

```

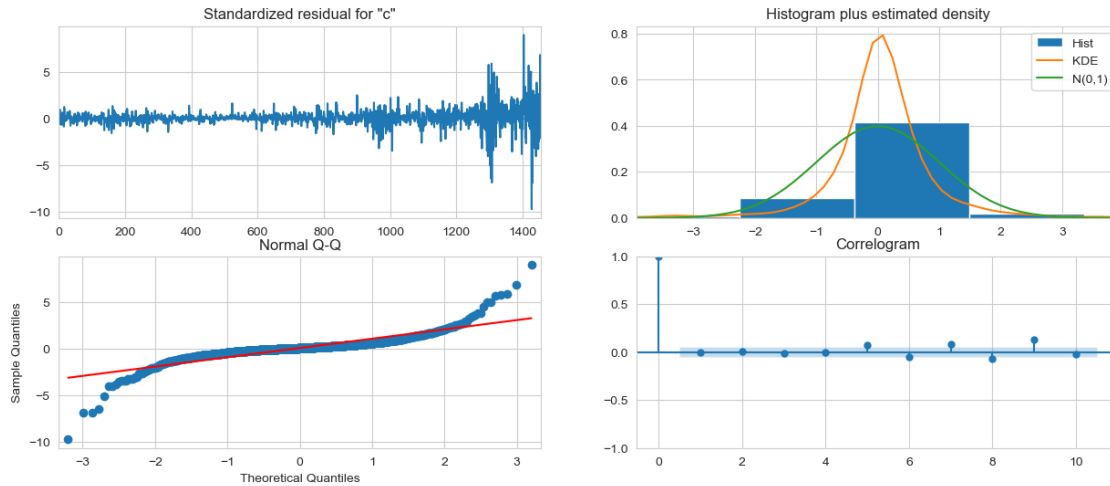
[54]: {'Mean Absolute Percentage Error (%)': 13.94,
      'Root Mean Squared Error': 27.427856712866763}

```


```

[55]: model.plot_diagnostics(figsize=(15,6))
      plt.show()

```



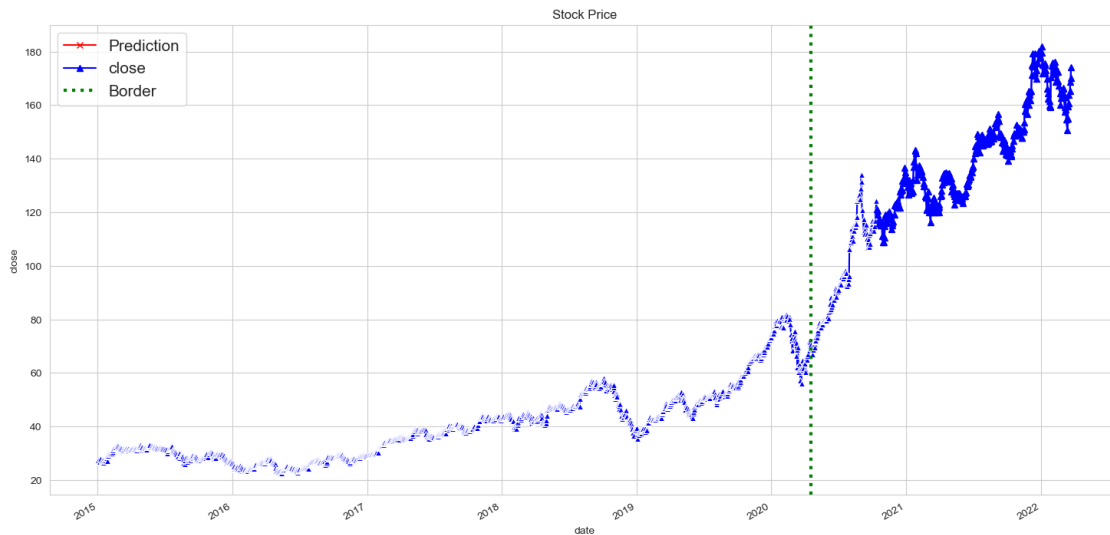
```
[56]: # plotting residuals and density
fig, ax = plt.subplots(1, 2)
fig.set_figwidth(16)
residuals = pd.DataFrame(model.resid)
residuals.plot(title="Residuals", ax=ax[0])
residuals.plot(kind="kde", title="Density", ax=ax[1])
plt.show()
```



```
[57]: price_forecast = model.forecast(12)
price_forecast = pd.Series(price_forecast, index=fy_forecast)
price_forecast = price_forecast.rename("Prediction")

fig, ax = plt.subplots(figsize=(18, 9))
sns.lineplot(x="date", y="close", data=df_close, color="blue", marker="^")
price_forecast.plot(ax=ax, c="red", marker="x", label="Prediction")
X_test.plot(ax=ax, c="blue", marker="^")
plt.title("Stock Price")
```

```
ax.axvline(x=18370, ls=":", linewidth=3, c="green", label="Border")
plt.legend(loc=0, fontsize=15)
plt.show()
```



```
[58]: # printing forecast predictions for next twelve months
df = pd.DataFrame({"Predictions": model.forecast(12)})
df["date"] = fy_forecast
df.set_index("date")
```

```
[58]:
```

	Predictions
date	
2022-04-01	121.747739
2022-05-01	121.507763
2022-06-01	121.596670
2022-07-01	121.563732
2022-08-01	121.575935
2022-09-01	121.571414
2022-10-01	121.573088
2022-11-01	121.572468
2022-12-01	121.572698
2023-01-01	121.572613
2023-02-01	121.572644
2023-03-01	121.572633

```
[59]: # splitting data into training and testing for an LSTM model
x_train, y_train = X_train.index, X_train.close
x_test, y_test = X_test.index, X_test.close

print(x_train.shape), print(x_test.shape)
```

```
(1456,)
(364,)
```

```
[59]: (None, None)
```

```
[60]: model = Sequential()
model.add(LSTM(50, return_sequences=True, input_shape=(100, 1)))
model.add(LSTM(50, return_sequences=True))
model.add(LSTM(50))
model.add(Dense(1))

model.compile(loss="mean_squared_error", optimizer="adam")
model.summary()
```

```
WARNING:tensorflow:From C:\Users\vince\anaconda3\Lib\site-
packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated.
Please use tf.compat.v1.get_default_graph instead.
```

```
WARNING:tensorflow:From C:\Users\vince\anaconda3\Lib\site-
packages\keras\src\optimizers\_init_.py:309: The name tf.train.Optimizer is
deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 100, 50)	10400
lstm_1 (LSTM)	(None, 100, 50)	20200
lstm_2 (LSTM)	(None, 50)	20200
dense (Dense)	(None, 1)	51

```
=====  
Total params: 50851 (198.64 KB)  
Trainable params: 50851 (198.64 KB)  
Non-trainable params: 0 (0.00 Byte)
```

```
-----  
model = Sequential(): Creates a new sequential model, which is a linear stack of layers.
```

```
model.add(LSTM(50, return_sequences=True, input_shape=(100, 1))): Adds the first LSTM layer
with 50 units (neurons), return_sequences=True indicates that the layer should return sequences
rather than a single output, and input_shape=(100, 1) specifies the input shape for the first layer
(100 time steps with 1 feature).
```

```
model.add(LSTM(50, return_sequences=True)): Adds a second LSTM layer with 50 units and
return_sequences=True, which means it also returns sequences.
```

model.add(LSTM(50)): Adds a third LSTM layer with 50 units, but return_sequences is not specified, so it defaults to False, meaning this layer will return a single output for the whole sequence.

model.add(Dense(1)): Adds a dense (fully connected) layer with 1 unit, which is the output layer for the model.

model.compile(loss="mean_squared_error", optimizer="adam"): Compiles the model with a mean squared error loss function and the Adam optimizer.

model.summary(): Prints a summary of the model, showing the architecture and the number of parameters in each layer.

```
[61]: # fitting the model on the testing data
model.fit(
    X_train,
    y_train,
    validation_data=(X_test, y_test),
    epochs=100,
    batch_size=64,
    verbose=1,
)
```

Epoch 1/100

WARNING:tensorflow:From C:\Users\vince\anaconda3\Lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

23/23 [=====] - 15s 148ms/step - loss: 2488.9712 - val_loss: 20388.2852

Epoch 2/100

23/23 [=====] - 1s 22ms/step - loss: 2381.2502 - val_loss: 19654.4473

Epoch 3/100

23/23 [=====] - 0s 14ms/step - loss: 2043.4578 - val_loss: 18389.5742

Epoch 4/100

23/23 [=====] - 0s 12ms/step - loss: 1755.6875 - val_loss: 17672.7148

Epoch 5/100

23/23 [=====] - 1s 24ms/step - loss: 1612.3430 - val_loss: 17255.2969

Epoch 6/100

23/23 [=====] - 1s 23ms/step - loss: 1518.8396 - val_loss: 16935.4551

Epoch 7/100

23/23 [=====] - 0s 14ms/step - loss: 1443.6001 - val_loss: 16650.6055

Epoch 8/100

23/23 [=====] - 0s 15ms/step - loss: 1377.0129 -
 val_loss: 16385.7891
 Epoch 9/100
 23/23 [=====] - 0s 12ms/step - loss: 1316.1627 -
 val_loss: 16136.6631
 Epoch 10/100
 23/23 [=====] - 0s 12ms/step - loss: 1259.8538 -
 val_loss: 15901.1621
 Epoch 11/100
 23/23 [=====] - 0s 12ms/step - loss: 1207.6207 -
 val_loss: 15675.2227
 Epoch 12/100
 23/23 [=====] - 0s 13ms/step - loss: 1158.9547 -
 val_loss: 15456.3301
 Epoch 13/100
 23/23 [=====] - 0s 12ms/step - loss: 1112.8175 -
 val_loss: 15247.6523
 Epoch 14/100
 23/23 [=====] - 0s 13ms/step - loss: 1070.1613 -
 val_loss: 15041.8857
 Epoch 15/100
 23/23 [=====] - 0s 13ms/step - loss: 1029.0826 -
 val_loss: 14847.8857
 Epoch 16/100
 23/23 [=====] - 0s 13ms/step - loss: 991.2211 -
 val_loss: 14656.4824
 Epoch 17/100
 23/23 [=====] - 0s 13ms/step - loss: 954.9692 -
 val_loss: 14472.6143
 Epoch 18/100
 23/23 [=====] - 0s 12ms/step - loss: 921.1567 -
 val_loss: 14291.8037
 Epoch 19/100
 23/23 [=====] - 0s 12ms/step - loss: 889.0429 -
 val_loss: 14117.3115
 Epoch 20/100
 23/23 [=====] - 0s 15ms/step - loss: 858.5605 -
 val_loss: 13950.3721
 Epoch 21/100
 23/23 [=====] - 0s 12ms/step - loss: 830.1619 -
 val_loss: 13786.1104
 Epoch 22/100
 23/23 [=====] - 0s 12ms/step - loss: 803.4775 -
 val_loss: 13625.8311
 Epoch 23/100
 23/23 [=====] - 0s 13ms/step - loss: 777.9852 -
 val_loss: 13472.5938
 Epoch 24/100

23/23 [=====] - 0s 13ms/step - loss: 754.1510 -
 val_loss: 13323.9248
 Epoch 25/100
 23/23 [=====] - 0s 13ms/step - loss: 731.8716 -
 val_loss: 13177.8086
 Epoch 26/100
 23/23 [=====] - 0s 12ms/step - loss: 710.6928 -
 val_loss: 13036.5957
 Epoch 27/100
 23/23 [=====] - 0s 17ms/step - loss: 690.9562 -
 val_loss: 12899.5293
 Epoch 28/100
 23/23 [=====] - 0s 14ms/step - loss: 672.3544 -
 val_loss: 12766.9521
 Epoch 29/100
 23/23 [=====] - 0s 13ms/step - loss: 654.9308 -
 val_loss: 12638.7168
 Epoch 30/100
 23/23 [=====] - 1s 23ms/step - loss: 638.6453 -
 val_loss: 12513.7861
 Epoch 31/100
 23/23 [=====] - 0s 16ms/step - loss: 623.3749 -
 val_loss: 12392.7334
 Epoch 32/100
 23/23 [=====] - 0s 13ms/step - loss: 609.1682 -
 val_loss: 12274.8926
 Epoch 33/100
 23/23 [=====] - 0s 13ms/step - loss: 595.8460 -
 val_loss: 12161.9893
 Epoch 34/100
 23/23 [=====] - 0s 12ms/step - loss: 583.5210 -
 val_loss: 12051.2197
 Epoch 35/100
 23/23 [=====] - 0s 12ms/step - loss: 572.0040 -
 val_loss: 11943.7871
 Epoch 36/100
 23/23 [=====] - 0s 13ms/step - loss: 561.0577 -
 val_loss: 11844.2129
 Epoch 37/100
 23/23 [=====] - 0s 16ms/step - loss: 551.3561 -
 val_loss: 11741.2607
 Epoch 38/100
 23/23 [=====] - 0s 13ms/step - loss: 541.8694 -
 val_loss: 11646.7129
 Epoch 39/100
 23/23 [=====] - 0s 15ms/step - loss: 533.4210 -
 val_loss: 11552.8994
 Epoch 40/100

23/23 [=====] - 0s 12ms/step - loss: 525.3023 -
 val_loss: 11466.8115
 Epoch 41/100
 23/23 [=====] - 0s 12ms/step - loss: 518.1283 -
 val_loss: 11377.9570
 Epoch 42/100
 23/23 [=====] - 0s 13ms/step - loss: 511.2739 -
 val_loss: 11294.5342
 Epoch 43/100
 23/23 [=====] - 0s 17ms/step - loss: 505.0655 -
 val_loss: 11213.3955
 Epoch 44/100
 23/23 [=====] - 0s 17ms/step - loss: 499.2321 -
 val_loss: 11138.4824
 Epoch 45/100
 23/23 [=====] - 0s 15ms/step - loss: 494.0583 -
 val_loss: 11063.6357
 Epoch 46/100
 23/23 [=====] - 0s 13ms/step - loss: 489.1286 -
 val_loss: 10994.4023
 Epoch 47/100
 23/23 [=====] - 0s 15ms/step - loss: 484.7482 -
 val_loss: 10925.1025
 Epoch 48/100
 23/23 [=====] - 0s 12ms/step - loss: 480.6475 -
 val_loss: 10859.4590
 Epoch 49/100
 23/23 [=====] - 0s 13ms/step - loss: 476.9771 -
 val_loss: 10793.3574
 Epoch 50/100
 23/23 [=====] - 0s 16ms/step - loss: 473.5978 -
 val_loss: 10731.1602
 Epoch 51/100
 23/23 [=====] - 0s 17ms/step - loss: 470.3553 -
 val_loss: 10679.4883
 Epoch 52/100
 23/23 [=====] - 0s 15ms/step - loss: 467.6563 -
 val_loss: 10622.9922
 Epoch 53/100
 23/23 [=====] - 0s 12ms/step - loss: 465.1718 -
 val_loss: 10566.1650
 Epoch 54/100
 23/23 [=====] - 0s 12ms/step - loss: 462.7691 -
 val_loss: 10516.8828
 Epoch 55/100
 23/23 [=====] - 0s 12ms/step - loss: 460.7469 -
 val_loss: 10467.3359
 Epoch 56/100

23/23 [=====] - 0s 13ms/step - loss: 458.9226 -
 val_loss: 10418.1943
 Epoch 57/100
 23/23 [=====] - 0s 14ms/step - loss: 457.1200 -
 val_loss: 10378.5713
 Epoch 58/100
 23/23 [=====] - 0s 13ms/step - loss: 455.6925 -
 val_loss: 10334.2988
 Epoch 59/100
 23/23 [=====] - 0s 12ms/step - loss: 454.2611 -
 val_loss: 10295.1631
 Epoch 60/100
 23/23 [=====] - 0s 13ms/step - loss: 453.0477 -
 val_loss: 10255.7773
 Epoch 61/100
 23/23 [=====] - 0s 12ms/step - loss: 451.9084 -
 val_loss: 10221.2090
 Epoch 62/100
 23/23 [=====] - 0s 13ms/step - loss: 450.9286 -
 val_loss: 10187.3184
 Epoch 63/100
 23/23 [=====] - 0s 15ms/step - loss: 449.9744 -
 val_loss: 10158.1250
 Epoch 64/100
 23/23 [=====] - 0s 16ms/step - loss: 437.2979 -
 val_loss: 10145.8506
 Epoch 65/100
 23/23 [=====] - 0s 12ms/step - loss: 367.6434 -
 val_loss: 10047.4893
 Epoch 66/100
 23/23 [=====] - 0s 14ms/step - loss: 354.7472 -
 val_loss: 9940.2900
 Epoch 67/100
 23/23 [=====] - 0s 15ms/step - loss: 343.8152 -
 val_loss: 9829.5586
 Epoch 68/100
 23/23 [=====] - 0s 15ms/step - loss: 334.7818 -
 val_loss: 9720.6641
 Epoch 69/100
 23/23 [=====] - 0s 15ms/step - loss: 325.9932 -
 val_loss: 9618.6777
 Epoch 70/100
 23/23 [=====] - 0s 15ms/step - loss: 317.7364 -
 val_loss: 9517.3125
 Epoch 71/100
 23/23 [=====] - 0s 16ms/step - loss: 309.6134 -
 val_loss: 9423.3691
 Epoch 72/100

23/23 [=====] - 0s 16ms/step - loss: 302.2291 -
val_loss: 9324.5146
Epoch 73/100
23/23 [=====] - 0s 12ms/step - loss: 294.8651 -
val_loss: 9232.8252
Epoch 74/100
23/23 [=====] - 0s 16ms/step - loss: 288.1949 -
val_loss: 9139.9346
Epoch 75/100
23/23 [=====] - 0s 18ms/step - loss: 281.7228 -
val_loss: 9048.2471
Epoch 76/100
23/23 [=====] - 0s 13ms/step - loss: 275.0583 -
val_loss: 8963.2119
Epoch 77/100
23/23 [=====] - 1s 22ms/step - loss: 269.1318 -
val_loss: 8870.8984
Epoch 78/100
23/23 [=====] - 0s 14ms/step - loss: 263.1180 -
val_loss: 8784.3125
Epoch 79/100
23/23 [=====] - 0s 15ms/step - loss: 257.3171 -
val_loss: 8700.8691
Epoch 80/100
23/23 [=====] - 0s 14ms/step - loss: 251.8498 -
val_loss: 8617.4131
Epoch 81/100
23/23 [=====] - 1s 22ms/step - loss: 246.5084 -
val_loss: 8535.0449
Epoch 82/100
23/23 [=====] - 0s 21ms/step - loss: 241.2217 -
val_loss: 8458.7432
Epoch 83/100
23/23 [=====] - 0s 12ms/step - loss: 236.2263 -
val_loss: 8377.5000
Epoch 84/100
23/23 [=====] - 0s 12ms/step - loss: 231.2125 -
val_loss: 8298.8252
Epoch 85/100
23/23 [=====] - 0s 12ms/step - loss: 226.4357 -
val_loss: 8221.0684
Epoch 86/100
23/23 [=====] - 0s 13ms/step - loss: 221.7800 -
val_loss: 8145.6196
Epoch 87/100
23/23 [=====] - 0s 12ms/step - loss: 217.3391 -
val_loss: 8067.7920
Epoch 88/100

```

23/23 [=====] - 0s 13ms/step - loss: 212.9360 -
val_loss: 7993.5112
Epoch 89/100
23/23 [=====] - 0s 11ms/step - loss: 208.5522 -
val_loss: 7923.5181
Epoch 90/100
23/23 [=====] - 0s 16ms/step - loss: 204.4295 -
val_loss: 7850.1909
Epoch 91/100
23/23 [=====] - 0s 21ms/step - loss: 200.2717 -
val_loss: 7775.6313
Epoch 92/100
23/23 [=====] - 0s 14ms/step - loss: 196.1105 -
val_loss: 7704.5234
Epoch 93/100
23/23 [=====] - 0s 13ms/step - loss: 192.1878 -
val_loss: 7634.9185
Epoch 94/100
23/23 [=====] - 1s 22ms/step - loss: 188.4186 -
val_loss: 7565.0781
Epoch 95/100
23/23 [=====] - 0s 21ms/step - loss: 184.5862 -
val_loss: 7495.7261
Epoch 96/100
23/23 [=====] - 1s 22ms/step - loss: 180.8233 -
val_loss: 7431.9038
Epoch 97/100
23/23 [=====] - 0s 22ms/step - loss: 177.3417 -
val_loss: 7363.1421
Epoch 98/100
23/23 [=====] - 1s 23ms/step - loss: 173.8050 -
val_loss: 7295.7637
Epoch 99/100
23/23 [=====] - 0s 14ms/step - loss: 170.3596 -
val_loss: 7229.7388
Epoch 100/100
23/23 [=====] - 0s 15ms/step - loss: 166.9794 -
val_loss: 7164.0039

```

[61]: <keras.src.callbacks.History at 0x2396af23110>

```

[62]: # validating the model
train_predict = model.predict(X_train)
test_predict = model.predict(X_test)
math.sqrt(mean_squared_error(y_train, train_predict))

```

```

46/46 [=====] - 3s 4ms/step
12/12 [=====] - 0s 9ms/step

```

```
[62]: 12.850884128834732
```

```
[63]: train_predict = model.predict(X_train)
      test_predict = model.predict(X_test)

      train_rmse = math.sqrt(mean_squared_error(y_train, train_predict))
      test_rmse = math.sqrt(mean_squared_error(y_test, test_predict))

      print(f"Training RMSE: {train_rmse}")
      print(f"Test RMSE: {test_rmse}")
```

```
46/46 [=====] - 0s 4ms/step
12/12 [=====] - 0s 8ms/step
Training RMSE: 12.850884128834732
Test RMSE: 84.6404395343665
```

```
[64]: # print the first ten predicted values of training set
      train_predict[:10]
```

```
[64]: array([[27.319466],
             [26.542522],
             [26.545076],
             [26.922934],
             [27.95377 ],
             [27.983282],
             [27.299484],
             [27.541065],
             [27.436642],
             [26.687561]], dtype=float32)
```

```
[12]: import os
      output_file = r'C:\path\to\local\directory\Untitled Folder\capstone_project_vdt.
      ↪pdf'
      os.system(f'jupyter nbconvert --to pdf --allow-chromium-download --output_
      ↪{output_file} capstone_project_vdt.ipynb')
```

```
[12]: 1
```